

Detection of Malicious Executable Files Using Machine Learning Algorithms

Kushagra Gupta
2022263

1. Motivation

Identifying malicious software has become increasingly complex as cyber threats evolve. Modern malware often employs sophisticated techniques like polymorphic layers or rapid self-updates to evade traditional detection methods based on signatures or heuristics. Consequently, it's becoming more difficult to accurately classify malware versus benign software.

To address this challenge, we plan to leverage the Microsoft malware dataset to explore the effectiveness of various machine-learning techniques. By analysing various executables, we aim to train models that can reliably distinguish between malicious and safe files. This research endeavour holds promise for enhancing our ability to combat emerging cyber threats and safeguard digital environments.

2. References and Related Literature:

Malware detection using Linear SVM by Baigaltugs et al.

The study converts each executable into a vector using frequency counting and TFIDF methods. The classification of vectors is done using a linear SVM, and different kernels are tested. The data consists of 40 different classes.

Malware Detection using Machine Learning by Dragos, Gavrilit et al.

The study extracts features from binary files and converts them into a binary vector with 308 features. The training is done to minimise the number of false positives, with the number of benign files being much greater than the number of malware files. Various perceptron algorithms like One-Sided Perceptron, Kernelized One-Sided Perceptron and Cascade Classification are utilised for classification.

3. Methods and outcomes

Mudit Bansal
2022306

Our goal is to assess how well various machine learning algorithms can classify malware. We also wish to explore feature reduction and feature analysis methods such as PCA and LDA, along with different methods for feature extraction to find out which features are best for malware classification.

The three metrics that will be used are the number of features and testing accuracy.

4. Database Description

The Microsoft Malware Classification dataset will be chosen, containing 10,868 samples. Each sample will contain a .asm file, which will be the assembly dump obtained using the IDA tool. Each sample will also have a name, which will be a 20-character-long unique hash value, and an integral label describing which class the malware belongs to, of which there will be 9 types. The assembly dump will also contain a metadata dump.

Due to the skewed nature of the data, we have also created an oversampled dataset which uses the synthetic minority oversampling technique to equalise the number of data points for each of the classes. After applying the technique, the final dataset we obtained is 26478 data points.

The data is split into training and testing sets following an 80-20 random split. The training set is further split to get a validation set using an 80-20 split.




Fig 1. Original Dataset Class Distribution




Fig 2. TSNE Plot of the Original Dataset




Fig 3. Heatmap of covariance matrix showing linearly dependent variables

5. Methodology and Models

The same method is followed for each dataset obtained to train three models: Naive Bayes, Random Forests, and Support Vector Machines. Individual steps carried out for each are described below.

5.1. Naive Bayes

Both original and oversampled datasets are considered.

5.2. Random Forest

Both original and oversampled datasets are utilised when training the models. Grid Search CV and Random Search CV are used to find the hyperparameters, number of estimators, max depth and the criterion (log-loss and GINI).

The features with information gain less than a certain threshold were dropped. The performance obtained on the reduced feature set gave a lower generalisation error.

About three features were found to have multicollinearity, and therefore, they were dropped.

5.3 Support Vector Machine

Original and oversampled datasets were considered, and three features were dropped to give better performance and remove multicollinearity.

6. Results and Analysis

We will use the log-loss of predicted probabilities and the accuracy score to compare different models. The performance trends observed are Gaussian Naive Bayes < Support Vector Machine < Random Forest regarding accuracy.

6.1 Naive Bayes

In the case of naive Bayes, there was a negligible difference in the results of both datasets. There was no noticeable difference in training and CV losses, indicating low variance.

6.2 Random Forest

Random forests perform the best among the models tested, with > 0.98 accuracy scores, and consistently benefit from oversampled datasets.

During this process, we also tried to ascertain the difference in time taken and performance of the RandomCV and GridSearchCV modules for hyperparameter tuning. GridSearchCV, due to checking all possible hyperparameters provided in a given range, tends to take much more computation time. In comparison, RandomSearchCV gives the best hyperparameter sampled from distributions for each hyperparameter tested. This makes it much faster to obtain a "good enough" set of hyperparameters, especially in scenarios where quick deployment is essential.

With the nature of malware constantly evolving, quick deployment of models could become essential, especially in critical applications. As expected, the hyperparameters provided by RandomSearchCV perform worse than the ones provided by GridSearchCV. However, with the difference in the accuracy being less than 0.5% in the case of RandomForest, the tradeoff regarding quicker training time can be helpful in certain scenarios.

6.3 SVM

The results of the SVM classifier were just slightly lower than the RF results, indicating our dataset's inherent inseparability. Like RF, there is a slight

increase in performance when using the oversampled dataset. The best-performing kernel was the sigmoid kernel.

7. Conclusions

We have shown that machine learning models based on simple features extracted from executables can be used to classify different types of malware. We also tried to understand the effect of features on the variance of a machine-learning model. Finally, we have tried to analyse the benefits and tradeoffs of two hyperparameter testing frameworks. To build on this work, in cases of availability of live executables, more quantitative information, such as usage of the network stack, memory usage, number of calls to specific functions, files read, etc., can be obtained during runtime, which can greatly help in the quick analysis and deployment of models to be more proactive against malware detection and classification.

7.1 Contributions

The work was divided among the team members on the following basis:

- Data Analysis: Kushagra
- Random Search and Grid Search: Mudit
- Support Vector Machine: Kushagra
- Random Forest: Mudit
- Compilation of Results: Kushagra
- Presentation: Mudit

8. References

- a. [Link to the Repository](#)
- b. B. Sanjaa and E. Chuluun, "Malware detection using linear SVM," Ifostr, 2013, pp. 136-138, doi: 10.1109/IFOST.2013.6616872.
- c. D. Gavriluț, M. Cimpoeșu, D. Anton and L. Ciortuz, "Malware detection using machine learning," 2009 International Multiconference on Computer Science and Information Technology, 2009, pp. 735-741, doi: 10.1109/IMC- SIT.2009.5352759